

---

# Multi-Agent Systems

Nov-Dec 2023

---

Lecturer: Eric Pauwels

[eric.pauwels@cwi.nl](mailto:eric.pauwels@cwi.nl)

---

# Overview of Today's Lecture

1. **Course Organisation**
  2. Agents and Multi-Agent Systems
  3. Agent Types
  4. Environments
  5. Topics discussed in this course
-

# Practical Information

## ■ General:

- Period: Nov-Dec 2023
- 6 ECTS

## ■ Teachers:

- **Lecturer:** Eric Pauwels ([eric.pauwels@cwi.nl](mailto:eric.pauwels@cwi.nl))
- **TAs:** Andrzej Szczepura, Abdallah Al-Janabi, Yassin Ben Allal

## ■ Recommended prior knowledge:

- Elementary calculus and probability, basic programming
- Schaum Outline Series (pdf online)

---

# Contact Hours

- Period:

- Nov - Dec (7 weeks)

- Plenary lectures:

- Tuesday: 15.30 - 17.15
  - Thursday: 11.00 - 12.45

- Lab sessions:

- Thu: 13:30-15:15, 15:30 - 17:15 and 17:30 - 19.15

- Final exam:

- Tuesday, Dec 19 (18:30) Emergo !!
-

# Lab sessions & Homework assignments

- Week 2-6: Report on *homework assignment*
    - *Grade: Pass or Fail: (2 Fails = 0/1)*
    - *Deadline: next lab session;*
    - *Groups of 4 students: single report*
  - *Week 7 : Individual homework, graded (4/4)*
    - *Deadline: to be announced;*
  - Final exam: 5/5
  - Total score: 10/10
-

# Course Material

## ■ Recommended Books: .

- Y. Shoham & K. Leyton-Brown:  
MultiAgent Systems, Cambridge

Online: <http://www.masfoundations.org/mas.pdf>

- R. S. Sutton & A.G. Barto:  
Reinforcement Learning, MIT Press ([2nd edition](#))
  - <http://incompleteideas.net/book/bookdraft2017nov5.pdf>
- William Spaniel: Game Theory 101
  - (book + YouTube channel)

You only need to know what was covered during lectures!

# Course Material: Optional

N. Nisan, T. Roughgarden, E. Tardos, V. Vazirani:  
Algorithmic Game Theory. Cambridge UP.

Solid, mathematical. Advanced.

A. Dixit, B. Nalebuff: Thinking Strategically. Norton.  
Lots of context and background. Interesting and non-technical.

D. Fudenberg and J. Tirole: Game Theory. MIT Press  
Solid, mathematical. Advanced.

---

You only need to know what was covered during lectures!

---


# Overview of Today's Lecture

1. Course Organisation
  2. Agents and Multi-Agent Systems
  3. Agent Types
  4. Environments
  5. Topics discussed in this Course
-



# What is an Agent?

An *agent* is a computer system that is situated in some *environment*, and that is capable of *autonomous* action in this environment in order to meet its delegated *objectives*



- Note: autonomy is a *spectrum*!

# Properties of Agents


- Live in some *environment*
- *Observe* this environment
- Maintain *knowledge* about the environment
- Make *decisions* about what to do
- *Act* in the environment
- *Communicate* with other systems/agents
  - *Coordinate* with other agents (**cooperative setting**)
  - *Negotiate* with other agents (**non-cooperative setting**)

# Multi-Agent Systems, a Definition

- A **Multi-Agent** System is one that consists of a number of agents that *interact (with each other and the environment)*
- In general, agents will have *different goals (often conflicting!)*
- To successfully interact, they will have to *learn, cooperate, coordinate, and negotiate*

# Agents and Environment

## Multi-agent Systems (MAS)



# Motivations for studying MAS

## **Technological:**

- Growth of distributed, **networked computer systems**
  - (computers act more as **individuals** than parts)
- **Robustness:** no single point of failure
- **Scalable and flexible:**
  - adding new agents when needed
  - asynchronous, parallel processing
- **Development and reusability**
  - components developed independently (by specialists)

# Application: Robotics

- Robots as Physical Agents (Embodiment)
  - Internet of Things (IoT)
  - Swarms of drones,
  - Fleet of autonomous vehicles
  - Physical internet



# Motivations for studying MAS (2)

## Scientific:

- Models for **interactivity** in (human) **societies**,
    - e.g. economics, social sciences
  - Models for **emergence of cooperation**
    - **Coordination:** cooperation among **non-antagonistic** agents
    - **Negotiation:** coordination among **self-interested** agents
-

# Multiagent Systems

- Typical **scientific questions** addressed:
  - ❑ What actions should agents take to **optimize** their rewards/utility?
  - ❑ How can self-interested agents **learn** from interaction with the environment and other agents to further their goals?
  - ❑ How can autonomous agents **coordinate** their activities so as to cooperatively achieve goals?
  - ❑ How can **cooperation** emerge in societies of self-interested agents?



# MAS as Distributed AI (DAI)

- **AI** : Cognitive processes in individuals
  - Inspiration: neuro-science, behaviourism, ...
- **DAI**: Social processes in groups
  - Inspiration: social sciences, economics, ....

# MAS as Distributed AI (DAI)

## Basic question in DAI

- How and when should which agents interact (compete or collaborate) in order to achieve their design objectives?

## Approaches:

- **Bottom-up:** given specific capabilities of individual agents, what collective behaviour will emerge?
- **Top-down:** Search for specific group-level rules (e.g. conventions, norms, etc.) that successfully constrain or guide behaviours at individual level;

---

# Overview of Today's Lecture

1. Course Organisation
  2. Agents and Multi-Agent Systems
  3. **Agent Types**
  4. Environments
  5. Topics discussed in this Course
-

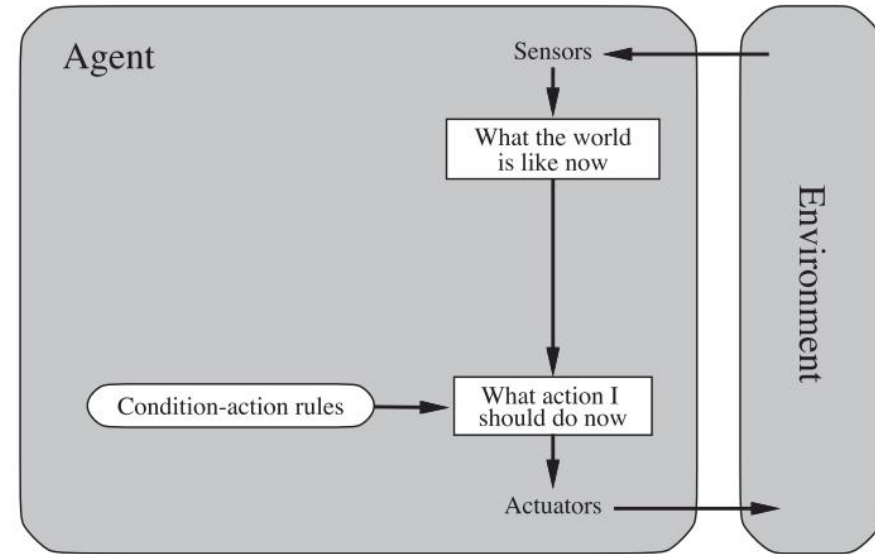
# Simple Typology for Agents

Intelligence in agents covers a spectrum:

- Reflex agents
  - Simple reflex agents
  - Model-based reflex agents
- Goal based agents
- Utility based agents
- Learning agents

# Type 1: Simple Reflex Agent

- Reacts to environment
  - Percept ----> Action  
Based on simple  
**if-then rules**  
(condition-action)
- Properties:
  - **No state**: ignore history
  - **Pre-computed** rules
  - **NO** Partial observability




---

## Type 2: Model-Based Reflex Agent

- Reflex agent with **state**
  - Agent uses memory to store an **internal representation** of its world
  - Internal **model based percept history**
  - This internal model allows him to **handle partially observable environment**
-

# Type 3: Goal-Based Agent

- **Goal** = desired outcome
- Goal-based (planning) agents act by **reasoning about which actions** to achieve the goal
- Less efficient, but more **adaptive** and **flexible**
- **Search and planning:** AI subfields concerned with **finding sequences of actions** to reach goal.



# Terminal vs. instrumental goals

- **Terminal goals:** you want them bcs... *you want them!* -- stop asking stupid questions :-)
- **Instrumental goals:** you want them bcs they *help* you *achieve your terminal goals*
- **Convergent instrumental goals:**  
instrumental to *wide range* of terminal goals:
  - ❑ Self preservation and self-improvement
  - ❑ Goal preservation
  - ❑ Resource acquisition



# Type 4: Utility-Based Agent

- **Utility**-based agents:
  - distinguish btw **goal and non-goal states**
- **Utility-based** agents use **utility function**
- **Utility function**
  - $U(\text{state})$  quantifies “**happiness**” (as real number)
  - **Preferred** world state has **higher utility** for agent
- **Strategic Agents:** Allows **rational decisions** in more situations
  - Evaluation of the **tradeoffs among conflicting** goals
  - Evaluation of **competing** goals

# Type 5: Learning Agent

## Four essential components

- **Actor:** responsible for selecting action in environment
- **Critic:** quantifies how well the agent is doing wrt. performance standard (e.g. utility).
- **Learner:** responsible for making improvements by learning from interactions.
- **Explorer:** responsible for suggesting new actions that will lead to novel and informative experiences.



# Intelligent Agents

- An *intelligent agent* is a computer system capable of *flexible* autonomous action in some environment
- **Autonomous:** not pre-determined by designer
- By *flexible*, we mean:
  1. *reactive*  
(able to receive information from environment and **respond**)
  2. *pro-active*  
(able to **reason and/or learn** and work towards **goals**)
  3. *social*  
(able to **communicate, coordinate, negotiate** and **cooperate**)


---

# Overview of Today's Lecture

1. Course Organisation
  2. Agents and Multi-Agent Systems
  3. Agent Types
  4. **Environments**
  5. Topics discussed in this Course
-


# Environments

- Agents act **in/on environments**




# Environments - *Accessible* vs. *inaccessible*

- An *accessible* environment: agent has **complete, accurate, up-to-date information** about the environment's state
- Most moderately complex environments (including, for example, the everyday physical world and the Internet) are *inaccessible*
- The **more accessible** an environment is, the **simpler** it is to build agents to operate in it



# Environments - *Deterministic* vs. *nondeterministic*

- A *deterministic* environment is one in which any **action** has a **single guaranteed effect** - there is **no uncertainty** about the state that will result from performing an action
- The physical world can to all intents and purposes be regarded as **non-deterministic**
- Non-deterministic environments present greater problems for the agent designer



# Environments - *Static* vs. *dynamic*

- A *static* environment is one that can be assumed to remain **unchanged except by the performance of actions** by the agent
- A dynamic environment is one that has other processes operating on it, and which hence **changes in ways beyond the agent's control**
- Other processes can interfere with the agent's actions (as in concurrent systems theory)





---

# Overview of Today's Lecture

1. Course Organisation
  2. Agents and Multi-Agent Systems
  3. Agent Types
  4. Environments
  5. Topics discussed in this Course
-

# Main topics in this course:

- **Strategic Agents:** Elementary **Game Theory:**
  - **Multiple** (stateless) **competing** agents,
  - **Rational choice** among multiple actions (max reward)
- **Exploration** versus **Exploitation**
- **Learning Agents:** **Reinforcement Learning**
  - **Single agent** learning to **optimize reward** from **sequential interactions with environment**;
  - **Multi-agent** Reinforcement Learning (MARL)
    - **Intro and some pointers**

# Summary

- An **agent** is a computer system that is capable of *autonomous* action in some *environment*, in order to achieve its delegated *objectives*
- Agent **properties**: reactive, proactive, social
- A **multi-agent system** is a system that consists of a number of agents, which *interact* with one *another* and the *environment*.
- Three **main topics** in this course:
  - **Strategic** agents (game theory)
  - **Exploration** versus **Exploitation**
  - **Learning** agents