

Multi-Agent Systems

Introduction to Reinforcement Learning

Part 4: Actor-Critic Algorithm

Eric Pauwels (CWI & VU)

December 14, 2023

Outline

Actor-Critic: Combining value- and policy-based learning

Advantage Actor-Critic (A2C)

Actor-Critic combines **valued-based** and **policy-based** learning.

Actor-Critic algorithms therefore have two components that are **learned jointly**:

- **Actor** learns a **parametrised policy**
- **Critic** learns **value function** to evaluate state-action pairs;

Advantage function: Select action based on how it **performs** relative to other actions in that state:

$$a_{\pi}(s, a) := q_{\pi}(s, a) - v_{\pi}(s)$$

Quick Recap

Policy gradient along trajectory τ

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

Restricting the trajectory to the part starting at s_t :

$$R_t(\tau) = R(s_t, a_t, \dots, s_T) \implies \mathbb{E}_{\tau \sim \pi_{\theta}} [R_t(\tau)] = q_{\pi_{\theta}}(s_t, a_t);$$


$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T q_{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

Policy Gradient via Monte Carlo Sampling

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right]$$

↓ single sample = single trajectory τ

$$\nabla_{\theta} J(\theta) \approx \sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$



$$R_t(\tau) = r_{t+1} + r_{t+2} + \dots + r_T \approx q_{\pi_{\theta}}(s_t, a_t)$$

↑
generated
 $s_t \xrightarrow{a_t} s_{t+1}$

↑
1-sample.

Policy gradient: Quick Recap

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T R_t(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] \quad (\text{MC}) \\ &= \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T q_{\pi_{\theta}}(s_t, a_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right] \quad (\text{value fct})\end{aligned}$$

For N sampled paths $\tau_i = \{s_{i,0}, a_{i,0}, s_{i,1}, r_{i,1}, \dots\}$:

$$\begin{aligned}\nabla_{\theta} J(\theta) &\approx \frac{1}{N} \sum_{i=1}^N \left[\sum_{t=0}^T R_t(\tau_i) \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \right] \\ &= \frac{1}{N} \sum_{i=1}^N \left[\sum_{t=0}^T q_{\pi_{\theta}}(s_{i,t}, a_{i,t}) \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \right]\end{aligned}$$

Actor-Critic Methods

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T \underbrace{\nabla_{\theta} \log \pi_{\theta}(a_t | s_t)}_{\text{ACTOR}} \underbrace{q_{\pi_{\theta}}(s_t, a_t)}_{\text{CRITIC}} \right]$$

- **Critic** estimates the **value function** (could be action-value q or state-value v function).
- **Actor** updates the **policy distribution** in the **direction suggested by the critic**, i.e.:
 - Changes the policy to **increase the likelihood of actions** that get **high values from the critic**. the critic, more l

Gradient Policy Theorem: Interpretation

$$\frac{dJ}{d\theta} = E_{\pi_{\theta}} \left[\sum_{t=0}^T q_{\pi_{\theta}}(s_t, a_t) \frac{d \log \pi_{\theta}(a_t | s_t)}{d\theta} \right]$$

Interpretation:


* If $q_{\pi_{\theta}}(s_t, a_t)$ "high", changing θ such that a_t becomes more likely (i.e.: $(\log) \pi_{\theta}(a_t | s_t) \uparrow$) increases $J(\theta)$

* Hence:

if $\frac{d \log \pi_{\theta}(a_t | s_t)}{d\theta} > 0 \Rightarrow \frac{dJ}{d\theta} > 0$
 increase θ to increase $J(\theta)$

if $\frac{d \log \pi_{\theta}(a_t | s_t)}{d\theta} < 0 \Rightarrow \frac{dJ}{d\theta} < 0$
 decrease θ to increase $J(\theta)$.

Gradient Policy Theorem: Interpretation




"Mixed message":

① $q(s, a_1) = 100 \rightarrow$ make a_1 more likely
 $\Rightarrow \theta \uparrow$

② $q(s, a_2) = 90 \rightarrow$ make a_2 more likely
 $\Rightarrow \theta \downarrow$

Gradient Policy Theorem: Interpretation

$$\frac{dJ}{d\theta} \approx \sum_{t=0}^T q_{\theta}(s_t, a_t) \frac{d}{d\theta} (\log \pi_{\theta}(a_t | s_t))$$




$$\frac{dJ}{d\theta} \approx 100 \cdot \text{positive} + 90 \cdot \text{negative} \approx 0$$

Gradient Policy Theorem: Introducing baselines

- “Raw” q -values don’t give clear feed-back signal
- **Baseline**: some average wrt which q -values can be compared:
- Comparing q -values to baseline results in clearer feedback!
- E.g.: baseline for state s is 95 ($b(s) = 95$):
 - $q(s, a_1) - b(s) = 100 - 95 = 5 > 0$
 - $q(s, a_2) - b(s) = 90 - 95 = -5 < 0$
- Recall: increasing θ makes a_1 (a_2) more (less) likely; hence:

$$\frac{d}{d\theta} \log \pi_{\theta}(a_1 | s) > 0 \quad (\text{positive}) \qquad \frac{d}{d\theta} \log \pi_{\theta}(a_2 | s) < 0 \quad (\text{negative})$$



Introducing baselines: A2C

- Policy gradient theorem:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) q(s_t, a_t) \right]$$

- **Problem:** value of $q(s, a)$ is not very informative;
- We need a reference point or **baseline**: **natural choice** = $v(s)$
- **Advantage:** Relative value of an action as compared to other actions in that state:

$$A(s, a) := q(s, a) - v(s)$$

- **Advantage actor-critic (A2C)**

$$\nabla_{\theta} J(\theta) \propto \mathbb{E}_{\tau} \left[\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) (q_{\pi_{\theta}}(s_t, a_t) - v_{\pi_{\theta}}(s_t)) \right]$$

Estimating Advantage

- Typically **two neural networks** to estimate
 - policy** $\rightarrow \pi_\theta$
 - value functions and advantage**: $\rightarrow v_w, q_w$
 - Network weights: θ and w
- Computational strategy**:
 - Estimate $v(s)$
 - Estimate $q(s, a)$ using Bellman eqs:

$$q(s_t, a_t) = \mathbb{E} [r_{t+1} + \gamma v(s_{t+1})]$$

- Along **sampled trajectory**:

$$\hat{q}(s_t, a_t) = r_{t+1} + \gamma \hat{v}(s_{t+1})$$

$$\hat{A}(s_t, a_t) = r_{t+1} + \gamma \hat{v}(s_{t+1}) - \hat{v}(s_t)$$

Estimating Advantage (2)


- n -step returns along sampled trajectory:

$$\hat{q}(s_t, a_t) = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n \hat{v}(s_{t+n})$$

- Combination of biased and unbiased estimate:
 - Actual returns: **unbiased but high variance**
(sample paths can be very different!)
 - **Bias** due to inclusion of estimate \hat{v} , **but lower variance**
(average over all actions);


Mathematical aside (1): Bias vs. Variance

- Consider sample $X_1, X_2, \dots, X_n \sim U(0, \theta)$ where θ is unknown and needs to be estimated;
- There are two natural estimators $\hat{\theta}$ for θ :
 - $\hat{\theta}_1 = \max_i X_i$: **biased** but **low variance**
 - $\hat{\theta}_2 = 2\bar{X} = 2/n \sum_i X_i$: **unbiased** but **high variance**



Mathematical aside (2a): Discount factor

Discount factor. $\gamma = \text{Prob of Continuing}$



$$\begin{aligned}
 ET &= \sum_{k=1}^{\infty} k \mathbb{P}(T=k) \\
 &= \sum_{k=1}^{\infty} k \cdot \gamma^k (1-\gamma) \\
 &= (1-\gamma) \gamma \underbrace{\sum_{k=1}^{\infty} k \gamma^{k-1}}
 \end{aligned}$$


Mathematical aside (2b): Discount factor

$$\begin{aligned}
 \sum_{k=1}^{\infty} k \gamma^{k-1} &= \frac{d}{d\gamma} \left(\sum_{k=1}^{\infty} \gamma^k \right) \\
 &= \frac{d}{d\gamma} \left(\frac{\gamma}{1-\gamma} \right) \\
 &= \frac{(1-\gamma) - \gamma(-1)}{(1-\gamma)^2} = \frac{1}{(1-\gamma)^2}
 \end{aligned}$$

$$ET = (1-\gamma) \gamma \frac{1}{(1-\gamma)^2} = \frac{\gamma}{1-\gamma}$$

Eg:
 $\gamma = 0.9$
 $ET \approx 9.$

Mathematical aside (2c): Discount factor



Further reading

[https://lilianweng.github.io/lil-log/2018/04/08/
policy-gradient-algorithms.html](https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html)